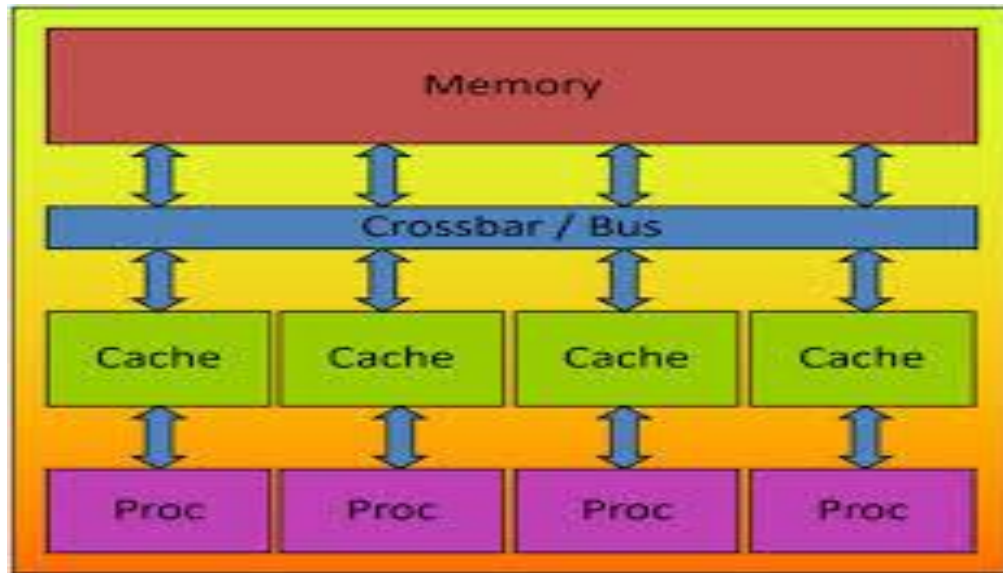


## Aim: Implement Min, Max, Sum and Average operations using Parallel Reduction



```
#include <iostream>
#include <vector>
#include <omp.h>
#include <climits>
using namespace std;
void min_reduction(vector<int>& arr) {
    int min_value = INT_MAX;
    #pragma omp parallel for reduction(min: min_value)
    for (int i = 0; i < arr.size(); i++) {
        if (arr[i] < min_value) {
            min_value = arr[i];
        }
    }
    cout << "Minimum value: " << min_value << endl;
}
void max_reduction(vector<int>& arr) {
    int max_value = INT_MIN;
    #pragma omp parallel for reduction(max: max_value)
    for (int i = 0; i < arr.size(); i++) {
        if (arr[i] > max_value) {
            max_value = arr[i];
        }
    }
}
```

```

}
}
cout << "Maximum value: " << max_value << endl;
}

void sum_reduction(vector<int>& arr) {

    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Sum: " << sum << endl;
}

void average_reduction(vector<int>& arr) {
    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Average: " << (double)sum / arr.size() << endl;
}

int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    min_reduction(arr);
    max_reduction(arr);
    sum_reduction(arr);
    average_reduction(arr);
}

```

Activities Terminal May 6 10:28 tanishka@tanishka-Lenovo: ~


```
#pragma omp parallel for reduction(max: max_value)
for (int i = 0; i < arr.size(); i++) {
    if (arr[i] > max_value) {
        max_value = arr[i];
    }
}
cout << "Maximum value: " << max_value << endl;
}

void sum_reduction(vector<int>& arr) {
    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Sum: " << sum << endl;
}

void average_reduction(vector<int>& arr) {
    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Average: " << (double)sum / arr.size() << endl;
}

int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    min_reduction(arr);
    max_reduction(arr);
    sum_reduction(arr);
    average_reduction(arr);
}
tanishka@tanishka-Lenovo:~$ g++ -o mms -fopenmp mms.cpp
tanishka@tanishka-Lenovo:~$ ./mms
Minimum value: 1
Maximum value: 9
Sum: 45
Average: 5
tanishka@tanishka-Lenovo:~$
```

**Aim: Write a program to implement Parallel Bubble Sort and Merge sort using OpenMP. Use existing algorithms and measure the performance of sequential and parallel algorithms.**



### Bubble Sort Code-

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
void bubble_sort_odd_even(vector<int>& arr) {
    bool isSorted = false;
    while (!isSorted) {
        isSorted = true;
        #pragma omp parallel for
        for (int i = 0; i < arr.size() - 1; i += 2) {
            if (arr[i] > arr[i + 1]) {
                swap(arr[i], arr[i + 1]);
                isSorted = false;
            }
        }
    }
    #pragma omp parallel for
```

```



for (int i = 1; i < arr.size() - 1; i += 2) {
    if (arr[i] > arr[i + 1]) {
        swap(arr[i], arr[i + 1]);
        isSorted = false;
    } } } }
int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    double start, end;
    // Measure performance of parallel bubble sort using odd-
    //even transposition
    start = omp_get_wtime();
    bubble_sort_odd_even(arr);
    end = omp_get_wtime();
    cout << "Parallel bubble sort using odd-even transposition time: " << end - start << endl;
}

```

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
void bubble_sort_odd_even(vector<int>& arr) {
    bool isSorted = false;
    while (!isSorted) {
        isSorted = true;
        #pragma omp parallel for
        for (int i = 0; i < arr.size() - 1; i += 2) {
            if (arr[i] > arr[i + 1]) {
                swap(arr[i], arr[i + 1]);
            }
            isSorted = false;
        }
        #pragma omp parallel for
        for (int i = 1; i < arr.size() - 1; i += 2) {
            if (arr[i] > arr[i + 1]) {
                swap(arr[i], arr[i + 1]);
            }
            isSorted = false;
        }
    }
}

int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    double start, end;
    // Measure performance of parallel bubble sort using odd-
    // even transposition
    start = omp_get_wtime();
    bubble_sort_odd_even(arr);
    end = omp_get_wtime();
    cout << "Parallel bubble sort using odd-even transposition time: " << end - start << endl;
}
tanishka@tanishka-Lenovo:~$ g++ -o bbsort -fopenmp bbsort.cpp
tanishka@tanishka-Lenovo:~$ ./bbsort
Parallel bubble sort using odd-even transposition time: 0.000706607
tanishka@tanishka-Lenovo:~$
```

# Merge Sort



```

#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
void merge(vector<int>& arr, int l, int m, int r) {
    int i, j, k;
    int n1 = m - l + 1;
    int n2 = r - m;
    vector<int> L(n1), R(n2);
    for (i = 0; i < n1; i++) {
        L[i] = arr[l + i];
    }
    for (j = 0; j < n2; j++) {
        R[j] = arr[m + 1 + j];
    }
    i = 0;
    j = 0;
    k = l;
    while (i < n1 && j < n2) {
        if (L[i] <= R[j]) {
            arr[k++] = L[i++];
        } else {
            arr[k++] = R[j++];
        }
    }
}
void merge_sort(vector<int>& arr, int l, int r) {
    if (l < r) {
        int m = l + (r - l) / 2;
#pragma omp task
        merge_sort(arr, l, m);
#pragma omp task
        merge_sort(arr, m + 1, r);
        merge(arr, l, m, r);
    }
}

```



```

    }
}
void parallel_merge_sort(vector<int>& arr) {
#pragma omp parallel
    {
#pragma omp single
        merge_sort(arr, 0, arr.size() - 1);
    }
}
int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    double start, end;
    // Measure performance of sequential merge sort
    start = omp_get_wtime();
    merge_sort(arr, 0, arr.size() - 1);
    end = omp_get_wtime();
    cout << "Sequential merge sort time: " << end - start << endl;
    // Measure performance of parallel merge sort
    arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    start = omp_get_wtime();
    parallel_merge_sort(arr);
    end = omp_get_wtime();
    return 0;
}

```



```
tanishka@tanishka-Lenovo:~$ g++ -o msort -fopenmp msort.cpp
```

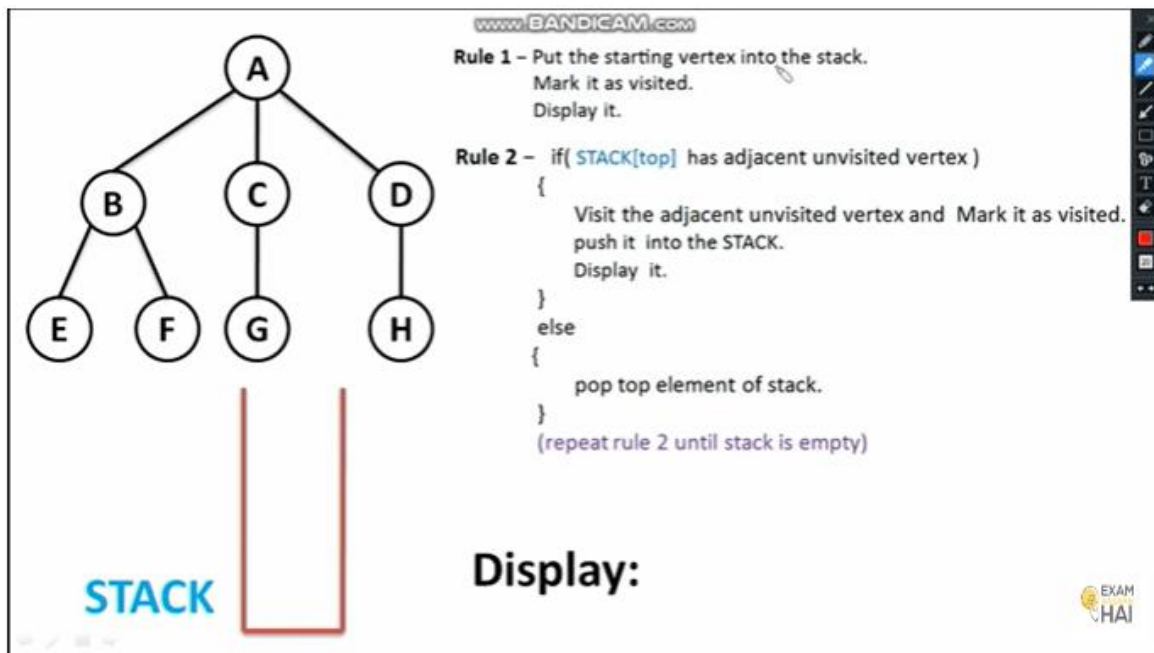
```
tanishka@tanishka-Lenovo:~$ ./msort
```

```
Sequential merge sort time: 3.1486e-05
```

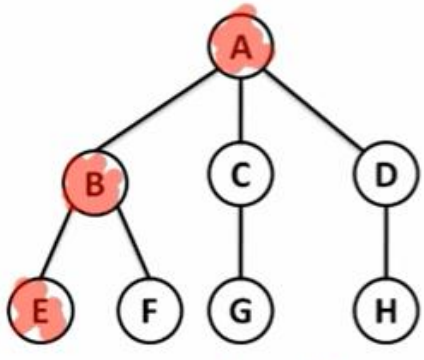
```
tanishka@tanishka-Lenovo:~$
```

**Aim: Design and implement Parallel Breadth First Search and Depth First Search based on existing algorithms using OpenMP. Use a Tree or an undirected graph for BFS and DFS.**

### DFS Concept -



www.BANDICAM.com



```

graph TD
    A((A)) --- B((B))
    A --- C((C))
    A --- D((D))
    B --- E((E))
    B --- F((F))
    C --- G((G))
    D --- H((H))
  
```

**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
Visit the adjacent unvisited vertex and Mark it as visited.  
push it into the STACK.  
Display it.  
}  
else  
{  
pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

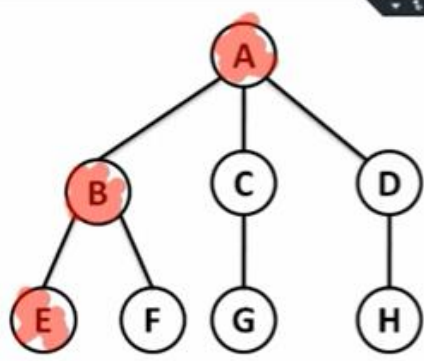
**STACK**: E, B, A

**Display**: ABE

9:31 / 17:40

Full screen (f)

www.BANDICAM.com



```

graph TD
    A((A)) --- B((B))
    A --- C((C))
    A --- D((D))
    B --- E((E))
    B --- F((F))
    C --- G((G))
    D --- H((H))
  
```

**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

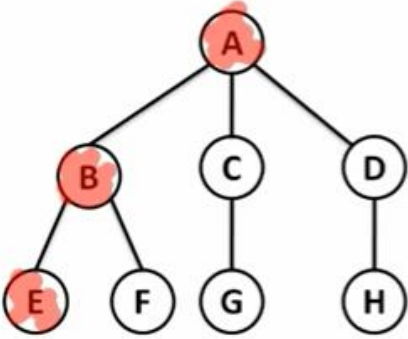
**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
Visit the adjacent unvisited vertex and Mark it as visited.  
push it into the STACK.  
Display it.  
}  
else  
{  
pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**STACK**: B, A

**Display**: ABE

10:05 / 17:40

www.BANDICAM.com



**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

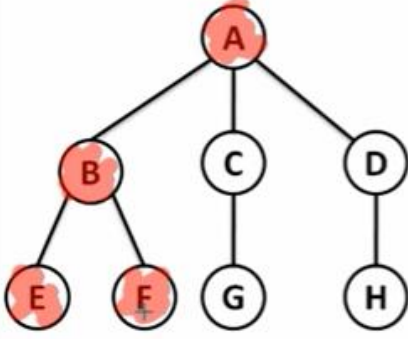
**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
    Visit the adjacent unvisited vertex and Mark it as visited.  
    push it into the **STACK**.  
    Display it.  
}  
else  
{  
    pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**STACK**

Display: AB

EXAM HAI

www.BANDICAM.com



**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

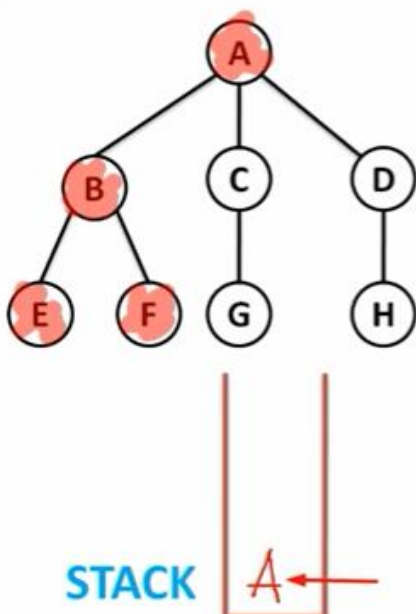
**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
    Visit the adjacent unvisited vertex and Mark it as visited.  
    push it into the **STACK**.  
    Display it.  
}  
else  
{  
    pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**STACK**

Display: AB EF

EXAM HAI

www.BANDICAM.com



**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

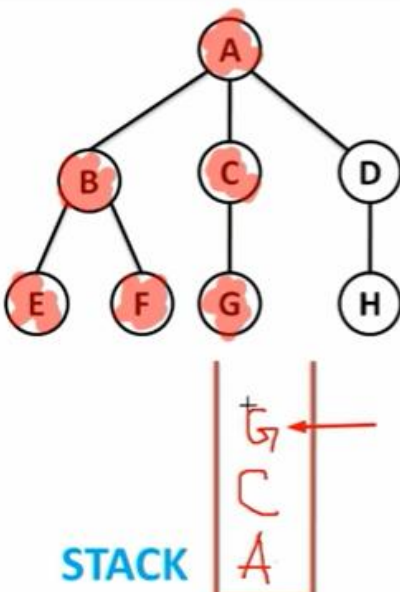
**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
    Visit the adjacent unvisited vertex and Mark it as visited.  
    push it into the **STACK**.  
    Display it.  
}  
else  
{  
    pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**STACK** A

**Display:** A B E F

EXAM HAI

www.BANDICAM.com



**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

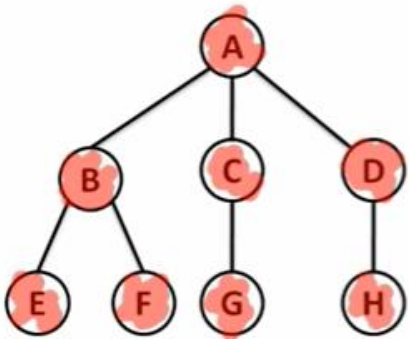
**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
    Visit the adjacent unvisited vertex and Mark it as visited.  
    push it into the **STACK**.  
    Display it.  
}  
else  
{  
    pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**STACK** G C A

**Display:** A B E F C G

EXAM HAI

www.BANDICAM.com



```

graph TD
    A((A)) --- B((B))
    A --- C((C))
    A --- D((D))
    B --- E((E))
    B --- F((F))
    C --- G((G))
    D --- H((H))
  
```

**STACK**

H  
D  
A

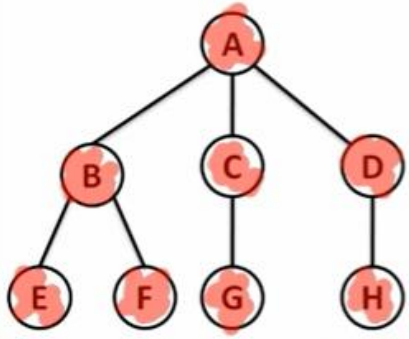
**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
Visit the adjacent unvisited vertex and Mark it as visited.  
push it into the **STACK**.  
Display it.  
}  
else  
{  
pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**Display:** ABEFCGDI

EXAM HAI

www.BANDICAM.com



```

graph TD
    A((A)) --- B((B))
    A --- C((C))
    A --- D((D))
    B --- E((E))
    B --- F((F))
    C --- G((G))
    D --- H((H))
  
```

**STACK**

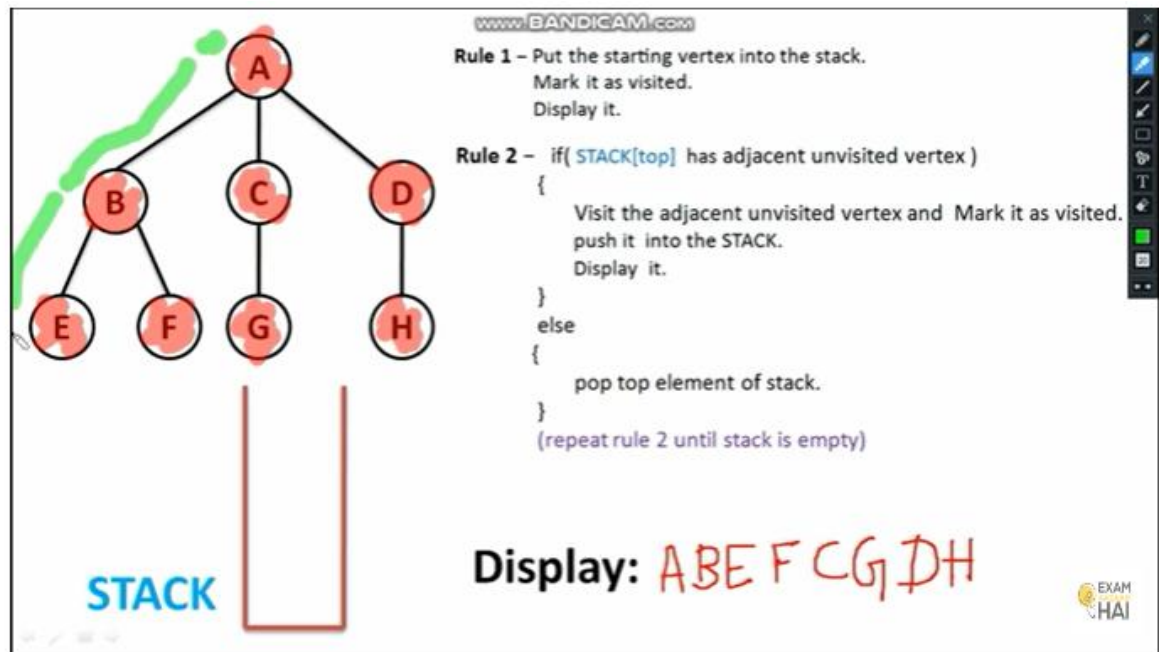
A

**Rule 1** – Put the starting vertex into the stack.  
Mark it as visited.  
Display it.

**Rule 2** – if( **STACK[top]** has adjacent unvisited vertex )  
{  
Visit the adjacent unvisited vertex and Mark it as visited.  
push it into the **STACK**.  
Display it.  
}  
else  
{  
pop top element of stack.  
}  
(repeat rule 2 until stack is empty)

**Display:** ABEFCGDH

EXAM HAI



### DFS Code-

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
const int MAXN = 1e5;
vector<int> adj[MAXN+5]; // adjacency list
bool visited[MAXN+5]; // mark visited nodes
void dfs(int node) {
    visited[node] = true;
    #pragma omp parallel for
    for (int i = 0; i < adj[node].size(); i++) {
        int next_node = adj[node][i];
        if (!visited[next_node]) {
            dfs(next_node);
        }
    }
}
int main() {
    cout << "Please enter nodes and edges";
    int n, m; // number of nodes and edges
    cin >> n >> m;
```



```
for (int i = 1; i <= m; i++) {
    int u, v; // edge between u and v
    cin >> u >> v;
    adj[u].push_back(v);
    adj[v].push_back(u);
}

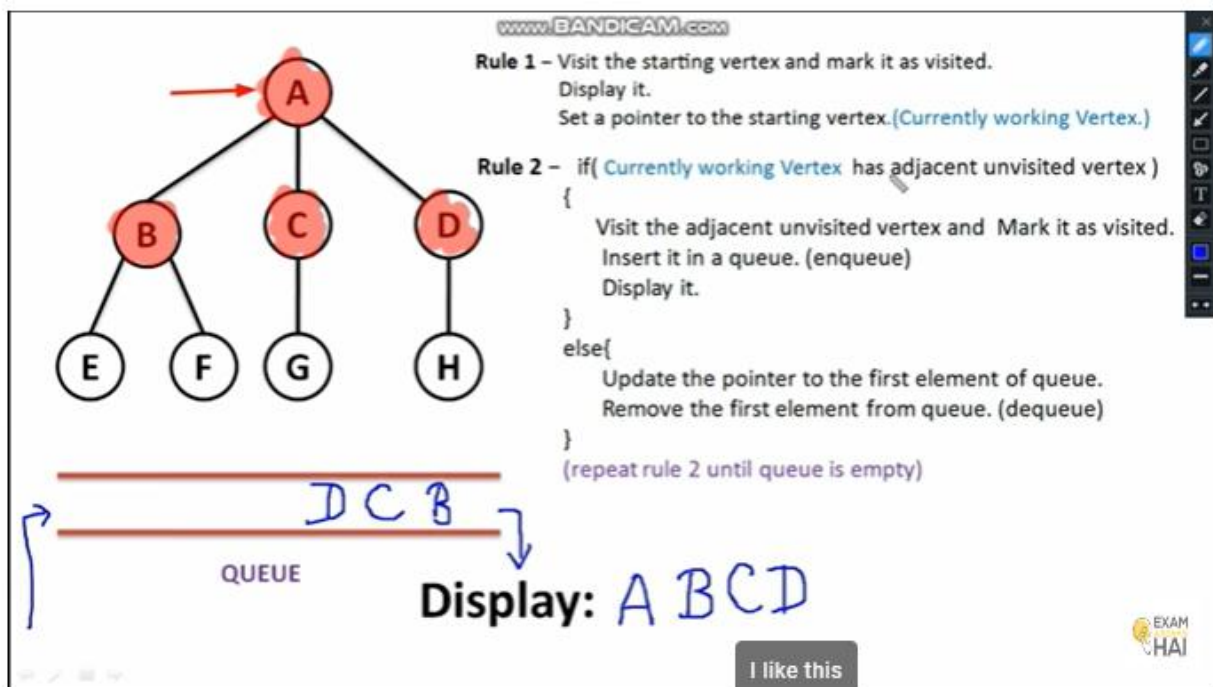
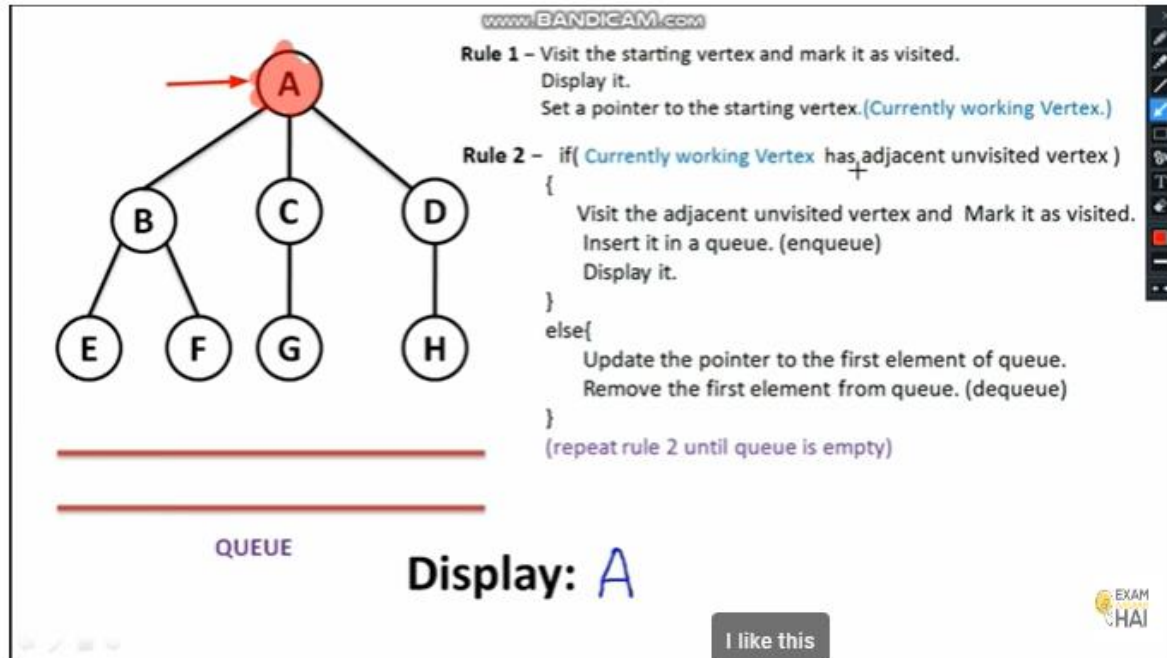
int start_node; // start node of DFS
cin >> start_node;

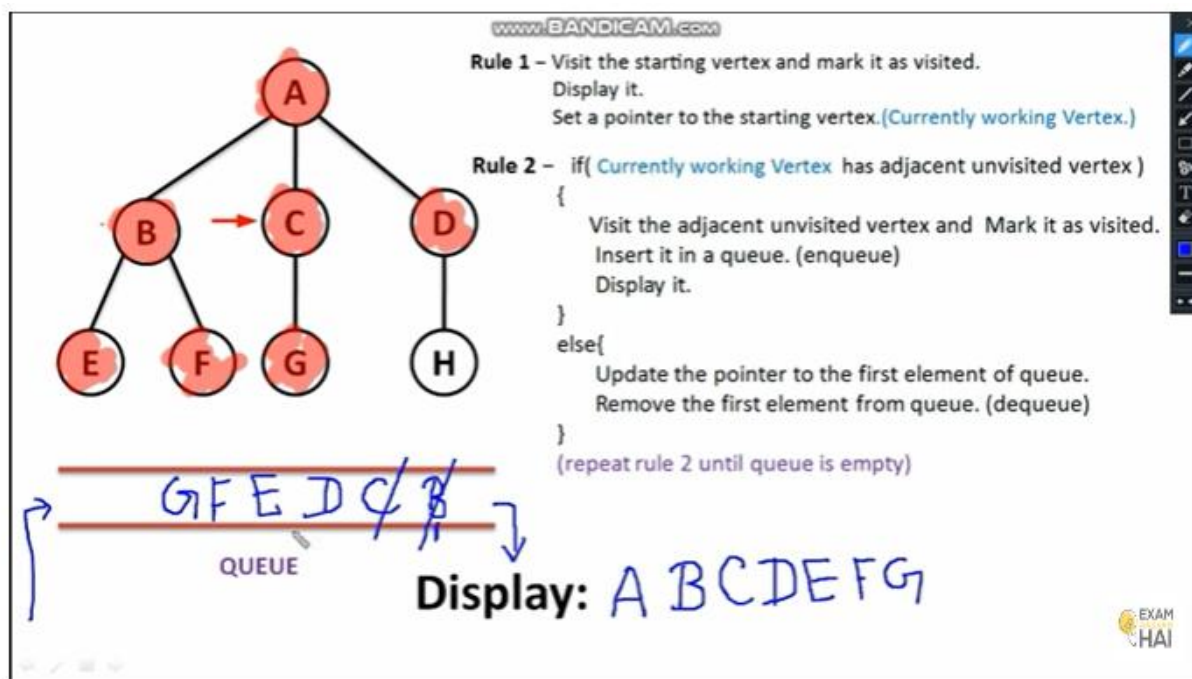
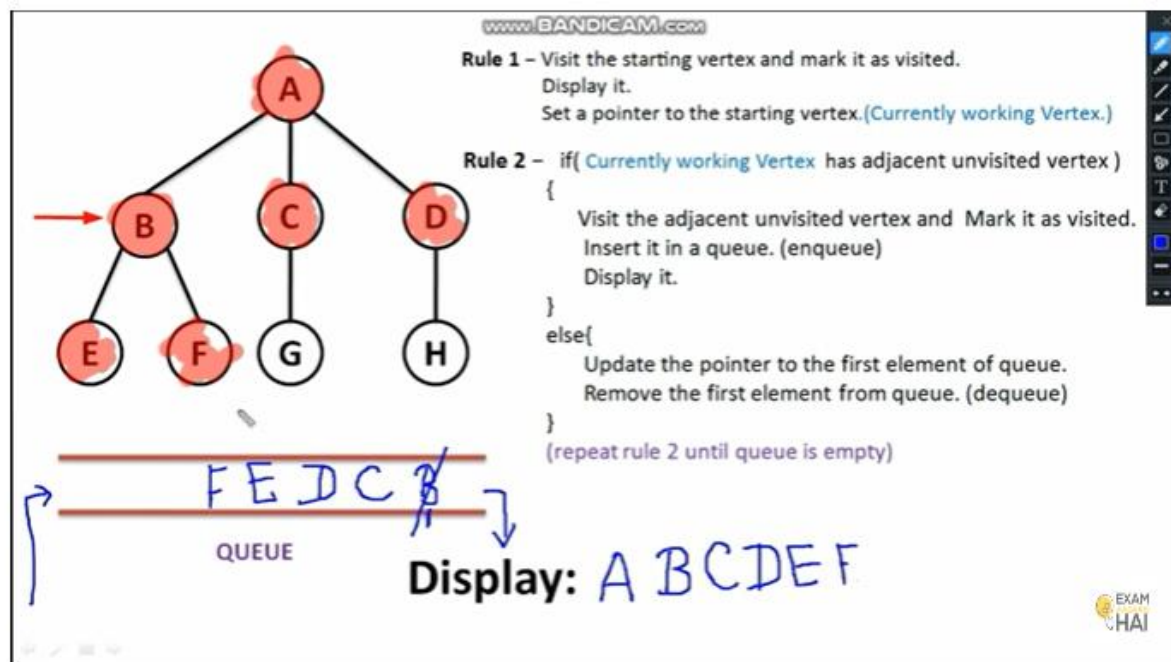
dfs(start_node);

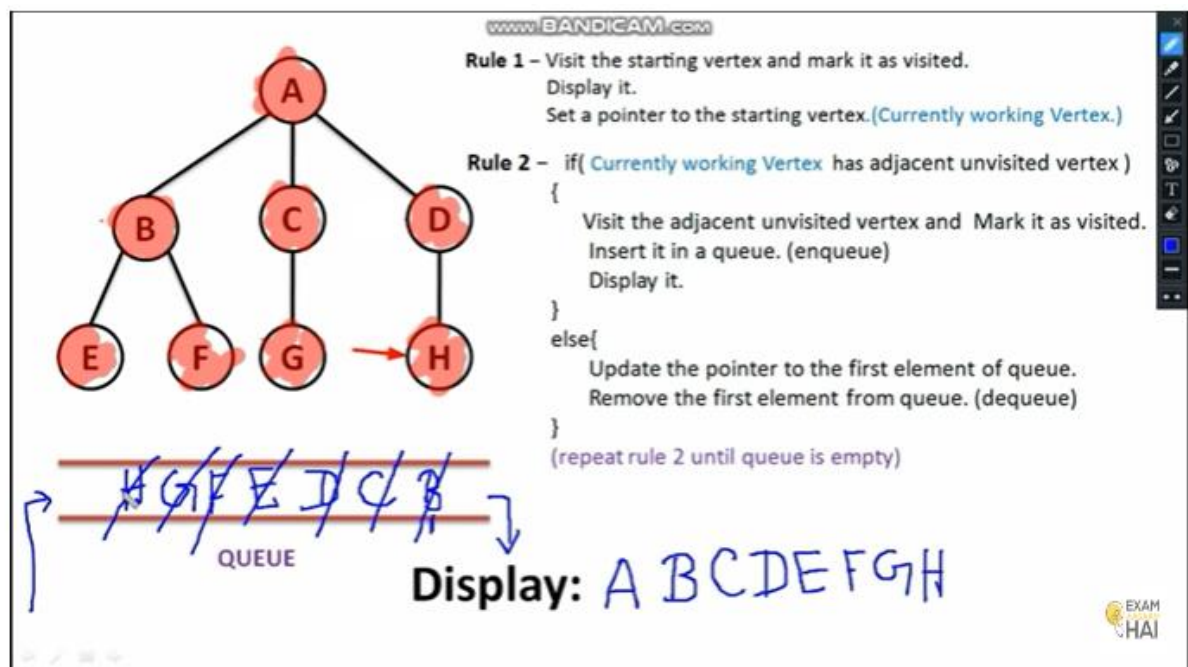
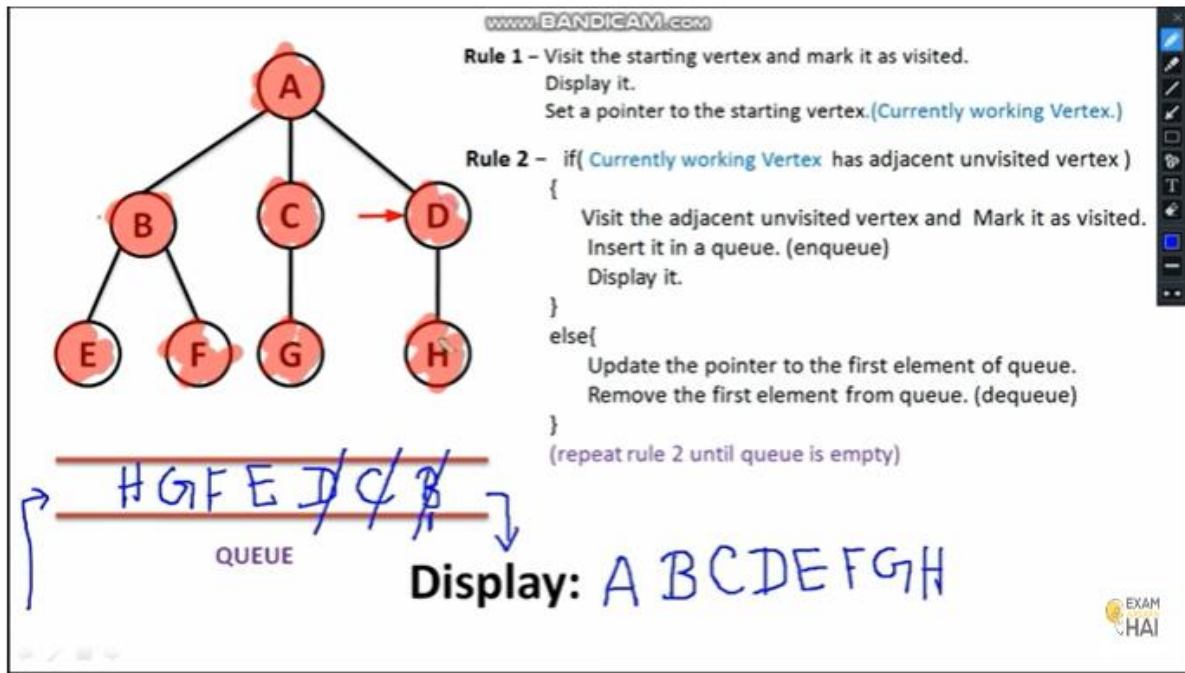
// Print visited nodes
for (int i = 1; i <= n; i++) {
    if (visited[i]) {
        cout << i << " ";
    }
}
cout << endl;

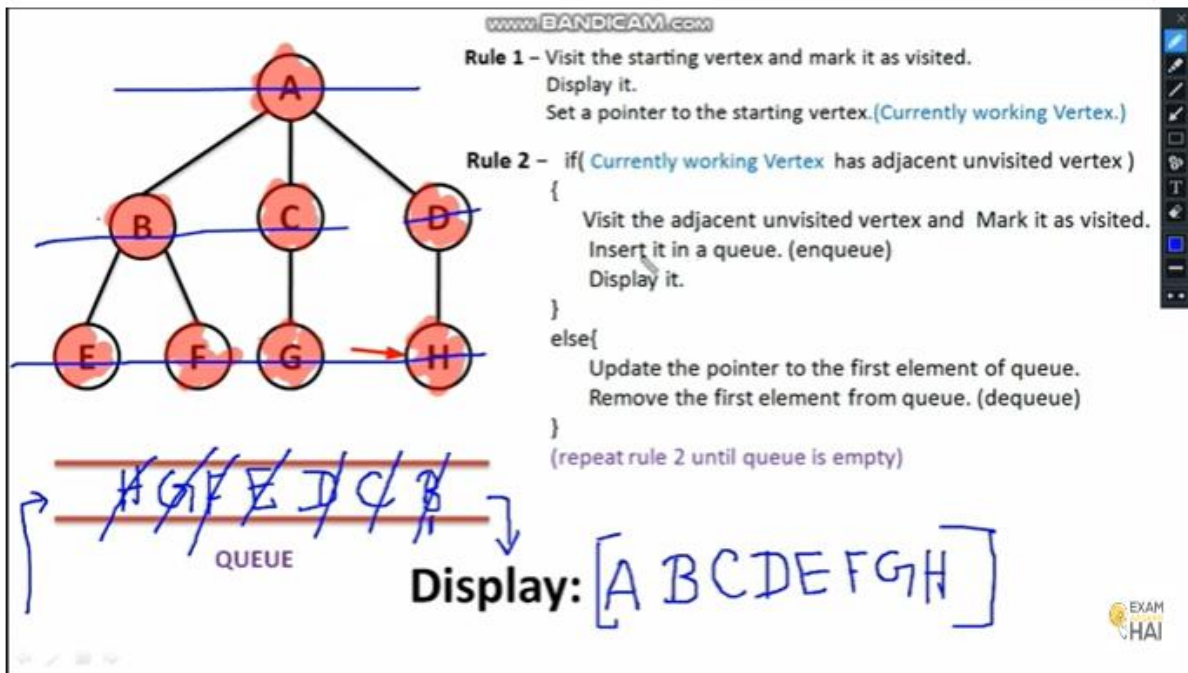
return 0;
}
```

## BFS Concept:









### BFS Code-

```
#include <iostream>
#include <queue>
#include <vector>
#include <omp.h>
```

```
using namespace std;
```

```
int main() {
    int num_vertices, num_edges, source;
    cin >> num_vertices >> num_edges >> source;

    vector<vector<int>> adj_list(num_vertices + 1);
    for (int i = 0; i < num_edges; i++) {
        int u, v;
        cin >> u >> v;
        adj_list[u].push_back(v);
        adj_list[v].push_back(u);
    }

    queue<int> q;
```

```

vector<bool> visited(num_vertices + 1, false);

q.push(source);
visited[source] = true;

while (!q.empty()) {
    int curr_vertex = q.front();
    q.pop();

    cout << curr_vertex << " ";

    #pragma omp parallel for shared(adj_list, visited, q) schedule(dynamic)
    for (int i = 0; i < adj_list[curr_vertex].size(); i++) {
        int neighbour = adj_list[curr_vertex][i];
        if (!visited[neighbour]) {
            visited[neighbour] = true;
            q.push(neighbour);
        }
    }
}

return 0;
}

```

```
Activities Terminal May 8 13:54 tanishka@tanishka-Lenovo: ~
tanishka@tanishka-Lenovo: ~
tanishka@tanishka-Lenovo:~$ cat >bfs8.cpp
#include <iostream>
#include <queue>
#include <vector>
#include <omp.h>
#include <cstring>

using namespace std;

const int MAXN = 100000;

vector<int> graph[MAXN+1];
int dist[MAXN+1];

void bfs(int start)
{
    queue<int> q;
    q.push(start);
    dist[start] = 0;

    while(!q.empty())
    {
        int u = q.front();
        q.pop();

        #pragma omp parallel for
        for(int i=0; i<graph[u].size(); i++)
        {
            int v = graph[u][i];
            if(dist[v] == -1)
            {
                dist[v] = dist[u] + 1;
                q.push(v);
            }
        }
    }
}
```

```
Activities Terminal May 8 13:54 tanishka@tanishka-Lenovo: ~
tanishka@tanishka-Lenovo: ~
tanishka@tanishka-Lenovo:~$ ./bfs9
6 8
0
0 1
0 2
1 2
2 0
2 3
3 3
1 4
4 5
0 2 0 3 tanishka@tanishka-Lenovo:~$ ./bfs9
5 7
0 1
0 2
1 2
2 0
2 3
3 3
4 4
0
0 4 0 1 3 2 tanishka@tanishka-Lenovo:~$ ./bfs9
4 4
1
1 2
2 4
1 3
3 4
1 3 0 4 tanishka@tanishka-Lenovo:~$ ./bfs9
4 4
0
0 1
1 3
0 2
2 3
0 2 1 3 tanishka@tanishka-Lenovo:~$
```

**Aim: Write a CUDA Program for :**

**1. Addition of two large vectors**

**2. Matrix Multiplication using CUDA C**

## **1. Vector Addition Program**

```
!nvcc --version
!pip install git+https://github.com/afnan47/cuda.git
%load_ext nvcc_plugin

%%cu
#include <iostream>
using namespace std;
__global__
void add(int* A, int* B, int* C, int size) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if (tid < size) {
        C[tid] = A[tid] + B[tid];
    }
}

void initialize(int* vector, int size) {
    for (int i = 0; i < size; i++) {
        vector[i] = rand() % 10;
    }
}

void print(int* vector, int size) {
    for (int i = 0; i < size; i++) {
        cout << vector[i] << " ";
    }
    cout << endl;
}

int main() {
    int N = 4;
```



```

int* A, * B, * C;
int vectorSize = N;
size_t vectorBytes = vectorSize * sizeof(int);
A = new int[vectorSize];
B = new int[vectorSize];
C = new int[vectorSize];
initialize(A, vectorSize);
initialize(B, vectorSize);
cout << "Vector A: ";
print(A, N);
cout << "Vector B: ";
print(B, N);
int* X, * Y, * Z;
cudaMalloc(&X, vectorBytes);
cudaMalloc(&Y, vectorBytes);
cudaMalloc(&Z, vectorBytes);
cudaMemcpy(X, A, vectorBytes, cudaMemcpyHostToDevice);
cudaMemcpy(Y, B, vectorBytes, cudaMemcpyHostToDevice);
int threadsPerBlock = 256;
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;
add<<<blocksPerGrid, threadsPerBlock>>>(X, Y, Z, N);
cudaMemcpy(C, Z, vectorBytes, cudaMemcpyDeviceToHost);
cout << "Addition: ";
print(C, N);
delete[] A;
delete[] B;
delete[] C;
cudaFree(X);
cudaFree(Y);
cudaFree(Z);
return 0;
}

```

Output:

Vector A: 3 6 7 5

Vector B: 3 5 6 2

Addition: 6 11 13 7

## 2. Matrix multiplication

```
!nvcc --version
```

```
!pip install git+https://github.com/afnan47/cuda.git
```

```
%load_ext nvcc_plugin
```

```
%%cu
```

```
#include <iostream>
```

```
#include <cuda.h>
```

```
using namespace std;
```

```
#define BLOCK_SIZE 2
```

```
__global__ void gpuMM(float *A, float *B, float *C, int N)
```

```
{
```

```
    // Matrix multiplication for NxN matrices C=A*B
```

```
    // Each thread computes a single element of C
```

```
    int row = blockIdx.y*blockDim.y + threadIdx.y;
```

```
    int col = blockIdx.x*blockDim.x + threadIdx.x;
```

```
    float sum = 0.f;
```

```

    for (int n = 0; n < N; ++n)
        sum += A[row*N+n]*B[n*N+col];

    C[row*N+col] = sum;
}

int main(int argc, char *argv[])
{
    int N; float K;

    // Perform matrix multiplication C = A*B
    // where A, B and C are NxN matrices
    // Restricted to matrices where N = K*BLOCK_SIZE;
    cout<<"Enter a Value for Size/2 of matrix";
    cin>>K;

    K = 1;
    N = K*BLOCK_SIZE;

    cout << "\n Executing Matrix Multiplication" << endl;
    cout << "\n Matrix size: " << N << "x" << N << endl;

    // Allocate memory on the host
    float *hA,*hB,*hC;
    hA = new float[N*N];
    hB = new float[N*N];
    hC = new float[N*N];

    // Initialize matrices on the host
    for (int j=0; j<N; j++){
        for (int i=0; i<N; i++){
            hA[j*N+i] = 2;
            hB[j*N+i] = 4;

        }
    }
}

```

```

// Allocate memory on the device
int size = N*N*sizeof(float); // Size of the memory in bytes
float *dA,*dB,*dC;
cudaMalloc(&dA,size);
cudaMalloc(&dB,size);
cudaMalloc(&dC,size);

dim3 threadBlock(BLOCK_SIZE,BLOCK_SIZE);
dim3 grid(K,K);
cout<<"\n Input Matrix 1 \n";
for (int row=0; row<N; row++){
    for (int col=0; col<N; col++){

        cout<<hA[row*col]<<" ";

    }
    cout<<endl;
}
cout<<"\n Input Matrix 2 \n";
for (int row=0; row<N; row++){
    for (int col=0; col<N; col++){

        cout<<hB[row*col]<<" ";

    }
    cout<<endl;
}

// Copy matrices from the host to device
cudaMemcpy(dA,hA,size,cudaMemcpyHostToDevice);
cudaMemcpy(dB,hB,size,cudaMemcpyHostToDevice);

//Execute the matrix multiplication kernel

gpuMM<<<grid,threadBlock>>>(dA,dB,dC,N);

```

```

// Now do the matrix multiplication on the CPU
/*float sum;
for (int row=0; row<N; row++){
    for (int col=0; col<N; col++){
        sum = 0.f;
        for (int n=0; n<N; n++){
            sum += hA[row*N+n]*hB[n*N+col];
        }
        hC[row*N+col] = sum;
        cout << sum << " ";

    }
    cout<<endl;
}*/

// Allocate memory to store the GPU answer on the host
float *C;
C = new float[N*N];

// Now copy the GPU result back to CPU
cudaMemcpy(C,dC,size,cudaMemcpyDeviceToHost);

// Check the result and make sure it is correct
cout << "\n\n\n\n Resultant matrix\n\n";
for (int row=0; row<N; row++){
    for (int col=0; col<N; col++){

        cout<<C[row*col]<< " ";

    }
    cout<<endl;
}

cout << "Finished." << endl;
}

```

Output-

Matrix size: 2x2

Input Matrix 1

2 2

2 2

Input Matrix 2

4 4

4 4

Resultant matrix

16 16

16 16

Finished.